

# Introduction

- To scale Elasticsearch, we need to add additional nodes eventually
- When doing so, replica shards are also assigned
- This approach cannot be used with Elastic Cloud
- We will be adding *two* nodes to our cluster, in two different ways
- The approaches are intended for development and *not* for production

# Approach #1 - the “right” way

1. Download & extract Elasticsearch
2. Configure the cluster and node names
3. Start up Elasticsearch

# Approach #2 - overwriting settings

1. Reuse an existing Elasticsearch directory
2. Start up Elasticsearch, overwriting the following settings at the command line
  - a. `cluster.name`
  - b. `node.name`
  - c. `path.data`
  - d. `path.logs`