

Introduction

- Prevent overwriting documents inadvertently due to concurrent operations
- There are many scenarios in which this can happen
 - E.g. handling concurrent visitors for a web application

How do we handle failures?

- Handle the situation at the application level
 - Retrieve the document again
 - Use `_primary_term` and `_seq_no` for a *new* update request
 - Remember to perform any calculations that use field values again

Lecture summary

- Sending write requests to Elasticsearch concurrently may overwrite changes made by other concurrent processes
- Traditionally, the `_version` field was used to prevent this
- Today, we use the `_primary_term` and `_seq_no` fields
- Elasticsearch will reject a write operation if it contains the wrong primary term or sequence number
 - This should be handled at the application level