Use explicit mappings

- Dynamic mapping is convenient, but often not a good idea in production

- Save disk space with optimized mappings when storing many documents

- Set `dynamic` to `"strict"`, not `false`

  - Avoids surprises and unexpected results

CodingExplained
CODE. SLEEP. REPEAT.

Mapping of text fields

- Don't always map strings as both `text` and `keyword`

  - Typically only one is needed

  - Each mapping requires disk space

- Do you need to perform full-text searches?

  - Add a `text` mapping

- Do you need to do aggregations, sorting, or filtering on exact values?

  - Add a `keyword` mapping

Disable coercion

- Coercion forgives you for not doing the right thing

- Try to do the right thing instead

- Always use the correct data types whenever possible

CodingExplained
CODE. SLEEP. REPEAT.

Use appropriate numeric data types

- For whole numbers, the `integer` data type might be enough

  - `long` can store larger numbers, but also uses more disk space

- For decimal numbers, the `float` data type might be precise enough

  - `double` stores numbers with a higher precision but uses 2x disk space

  - Usually, `float` provides enough precision

CodingExplained
CODE. SLEEP. REPEAT.

Mapping parameters

- Set `doc_values` to `false` if you don't need sorting, aggregations, and scripting

- Set `norms` to `false` if you don't need relevance scoring

- Set `index` to `false` if you don't need to filter on values

  - You can still do aggregations, e.g. for time series data

- Probably only worth the effort when storing lots of documents

  - Otherwise it's probably an over complication

- Worst case scenario, you will need to reindex documents